

## Assignment 2

Name: Fubang ZHAO

Date: 23/10/2017

### Question 1: Join Algorithms

Relation R contains 10,000 tuples and has 10 tuples per page. Relation S contains 4,000 tuples and also has 10 tuples per page. Attribute b of relation S is the primary key for S. Both relations are stored as simple heap files. Neither relation has any indexes built on it. 41 buffer pages are available.

- (1)  $L = \min(M, n) = 400$ . In this question, we have  $L/(B - 1) = 400/(41 - 1) \leq B - 2 = 39$ , so we get the cost:

$$3 \times (1000 + 400) = 4200I/Os \quad (1)$$

To remain the cost unchanged, we need to satisfy  $L/(B - 1) = B - 2$ , which means:

$$B > \sqrt{L} + 1 = 21 \quad (2)$$

So the minimum number of buffer pages is **22**.

- (2) The best way is to use Block Nested Loops. We set S as the outer, so the cost of Block Nested Loops is  $400 + \lceil 400/(B - 2) \rceil * 1000$ . Let **B = 402**, the cost is **1000 + 400 = 1400**
- (3) (a). Yes. Assuming that we partition R into k buckets, we need to satisfy:  $k - 1 + \lceil \frac{400}{k} \rceil + 1 + 1 \leq B$ , we get **k = 20**. And it satisfies that  $k < B - 2$ . So we have enough memory to perform the hybrid hash join.
- (b). For R, each bucket has  $\lceil \frac{400}{k} \rceil = 20$  pages. For S, each bucket has  $\lceil \frac{1000}{k} \rceil = 50$ . So there are  $20 + 50 = 70$  pages which have no need to write and reread. So the total I/O cost is:

$$3 \times (1000 + 400 - 70) + 70 = 4060I/Os \quad (3)$$

- (4) The idea is quite similar to the last question. However, we don't need a buffer to store S to do the merge. so we need to just satisfy  $k - 1 + \lceil \frac{400}{k} \rceil + 1 \leq B$ , which means  $k = 16$ . There are 25 pages which belong to  $S_0$  and have no need to write and reread. So the total cost is:

$$3 \times (400 - 25) + 25 = 1150I/Os \quad (4)$$

### Question 2

- (1) **Pass 0**: Find the best access path to  $R, S_1, S_2, \dots, S_{n-1}$  with  $n$  computations.

**Pass 1:** Find the best plans for all 2-relations sets. We need to spend  $2(n-1)$  computations because we need to consider each relation to another. In the end,  $n-1$  best plans are stored.

**Pass 2:** Find the best plan for all 3-relation sets with  $(n-1) * (n-2)$  computations and  $C_{n-1}^2$  plans are stored.

...

As shown in the table:

Pass Number	Nb of computations	Nb of plans stored
0	n	n
1	$2(n-1)$	$C_{n-1}^1$
2	$(n-2) * C_{n-1}^1$	$C_{n-1}^2$
...	...	...
k	$(n-k)C_{n-1}^{k-1} = kC_{n-1}^k$	$C_{n-1}^k$

So the total number of computations is:

$$n + 2(n-1) + \sum_{k=2}^{n-1} kC_{n-1}^k = 2n-1 + \sum_{k=1}^{n-1} kC_{n-1}^k = 2n-1 + (n-1)2^{n-2} \quad (5)$$

So the complexity of dynamic programming for finding an optimal plan is:  $O(n2^{n-2})$

- (2) As explained in last question, for Pass k, we need to store  $C_{n-1}^k$ . So the maximum number of plans is  $C_{n-1}^{\lfloor \frac{n-1}{2} \rfloor}$ .

### Question 3

- (1) As shown in the query plan:

```
-----
Bitmap Heap Scan on authors (cost=7918.93..22522.87 rows=321355 width=19)
(actual time=52.804..162.421 rows=322803 loops=1)
  Recheck Cond: ((name)::text < 'David J. DeWitt'::text)
  Heap Blocks: exact=10587
    -> Bitmap Index Scan on authors_name (cost=0.00..7838.59 rows=321355 width=0)
        (actual time=50.527..50.527 rows=322803 loops=1)
            Index Cond: ((name)::text < 'David J. DeWitt'::text)
Planning time: 0.090 ms
Execution time: 214.746 ms
(7 rows)
```

- (a) PostgreSQL uses the Bitmap Heap Scan on authors and Bitmap Index Scan on authors for this query.
- (b) The estimated number of rows is **321355** and the actual number of rows is **322803**.

(c) As shown in the picture below:

```
dblp=# select relname, reltuples from pg_class where relname='authors_name'
;
 relname      | reltuples 
-----+-----
 authors_name | 1.69134e+06
(1 row)
```

The estimation is based on the histogram and the estimated number of rows (tuples) before the bucket to which '*DavidJ.DeWitt*' belongs. As shown in the picture above, the total number of tuples is  $1.69134e + 06$ .

According to the *pg\_stats*, we can get '*DavidJ.DeWitt*' belongs to the 20th bucket out of 100 buckets. So the estimated Nb of rows is  $(1.69134e + 06) \times 19/100 \approx 321355$

(2) As shown in the query plan:

```
-----
Index Scan using authors_name on authors (cost=0.43..8.45 rows=1 width=19)
(actual time=0.025..0.026 rows=1 loops=1)
  Index Cond: ((name)::text = 'David J. DeWitt'::text)
Planning time: 1.856 ms
Execution time: 0.049 ms
(4 rows)
```

- (a) PostgreSQL uses **Index Scan** for this query.
  - (b) 1. This query has an equality condition and 'name' differs a lot so that there are not many rows.
  - 2. Compared to this query, the last question has a quite large number of rows(because the condition is '<', so it uses the bitmap index scan.
- (3) (a) As shown in the query plan:

```

Nested Loop (cost=8.89..163545.70 rows=5 width=159) (actual time=24.203..
4013.110 rows=186 loops=1)
  -> Hash Join (cost=8.46..163543.23 rows=5 width=27) (actual time=24.16
4..4007.708 rows=186 loops=1)
    Hash Cond: (pa.authid = a.id)
    -> Seq Scan on paperauths pa (cost=0.00..129792.71 rows=8997871
width=8) (actual time=0.017..1944.746 rows=8997871 loops=1)
    -> Hash (cost=8.45..8.45 rows=1 width=19) (actual time=0.016..0.
016 rows=1 loops=1)
      Buckets: 1024 Batches: 1 Memory Usage: 9kB
      -> Index Scan using authors_name on authors a (cost=0.43..
8.45 rows=1 width=19) (actual time=0.012..0.013 rows=1 loops=1)
        Index Cond: ((name)::text = 'David J. DeWitt'::text)
    -> Index Scan using papers_pkey on papers p (cost=0.43..0.48 rows=1 wi
dth=132) (actual time=0.024..0.024 rows=1 loops=186)
      Index Cond: (id = pa.paperid)
Planning time: 0.943 ms
Execution time: 4013.249 ms
(12 rows)

```

- (b) 1.  $a.name = 'DavidJ.DeWitt'$ . Implemented by the Index Scan using  $authors_{name}$  on authors a.  
 2.  $pa.authid = a.id$ . Implemented by the Hash Join which includes two operations: the Seq Scan on  $pa.authid$  and Hash on  $a.id$ .  
 3.  $pa.paperid = p.id$ . Implemented by the Nested Loop Join which includes the Hash Join and Index Scan above.
- (c) The estimated number of rows is **5** and the actual number of rows is **186**
- (d) Because that  $pa.authid$  is a foreign key and  $a.id$  is a primary key. So the estimated Nb of rows is:

$$N_{tuple_{pa}} \times \frac{1}{N_{tuple_a}} = \frac{8997871}{1.69134e+06} \approx 5 \quad (6)$$

- (4) (a) As shown in the query plan:

```

Hash Join (cost=225078.04..590553.72 rows=1709598 width=159) (actual time
=2866.383..15699.360 rows=1633612 loops=1)
  Hash Cond: (pa.paperid = p.id)
    -> Hash Join (cost=28422.80..281232.51 rows=1709598 width=27) (actual
time=271.311..7764.730 rows=1633612 loops=1)
      Hash Cond: (pa.authid = a.id)
        -> Seq Scan on paperauths pa (cost=0.00..129792.71 rows=8997871
width=8) (actual time=0.031..2093.764 rows=8997871 loops=1)
        -> Hash (cost=22522.87..22522.87 rows=321355 width=19) (actual t
ime=270.835..270.835 rows=322803 loops=1)
          Buckets: 65536 Batches: 8 Memory Usage: 2561kB
          -> Bitmap Heap Scan on authors a (cost=7918.93..22522.87 r
ows=321355 width=19) (actual time=53.042..162.779 rows=322803 loops=1)
            Recheck Cond: ((name)::text < 'David J. DeWitt'::text)
            Heap Blocks: exact=10587
            -> Bitmap Index Scan on authors_name (cost=0.00..783
8.59 rows=321355 width=0) (actual time=50.774..50.774 rows=322803 loops=1)
              Index Cond: ((name)::text < 'David J. DeWitt'::t
ext)
        -> Hash (cost=95777.44..95777.44 rows=3149344 width=132) (actual time=
2592.526..2592.526 rows=3150923 loops=1)
          Buckets: 32768 Batches: 256 Memory Usage: 2228kB
          -> Seq Scan on papers p (cost=0.00..95777.44 rows=3149344 width=
132) (actual time=0.004..925.594 rows=3150923 loops=1)
    Planning time: 0.562 ms
    Execution time: 15953.088 ms
(17 rows)

```

- (b) When we use Nested Loop Join, the cost is  $M + M * N$ . In the case of '=',  $M$  is very small, so it is efficient to use Nested Loop Join. However, in the case of '<',  $M$  is large so we choose to use Hash Join which cost is  $M + k * N$  ( $k$  depends on the size of memory, we can make it equal to 1 if we have enough memory).
- (c) For the Hash Join on  $pa.authid = a.id$ , the estimated number of rows is computed by:  $Size(Paperauths) \times RF_{a.name}$ , where:

$$RF_{a.name} = \frac{N_{estimated.rows}}{N_{tuples_a}} = \frac{321355}{1.69134e + 06} \quad (7)$$

$$Size(Paperauths) = 8997871 \quad (8)$$

So the estimated number of Rows is  $8997871 * \frac{321355}{1691340} \approx 1709615$